

Code & Coffee #22

Low-latency Java: breaking the boundaries

Ruslan Shevchenko
<ruslan@shevchenko.kiev.ua>
@rssh1

Overview

- ❖ What is low-latency and when it needed.
- ❖ LL-specific programming techniques:
 - ❖ Concurrent flows
 - ❖ Memory
 - ❖ Layout issues,
 - ❖ GC
 - ❖ Unmanaged memory access.
 - ❖ JNI

HPC != HPC



throughput



efficiency



speed



compromise

// low latency

low-latency java:

- ❖ Ineffective **// NEEDED: SOFT REALTIME**
- ❖ hardware is underused.
- ❖ Non-scalable
- ❖ operational range is strictly defined.
- ❖ Uncomfortable
- ❖ api reflect low-level details.

soft realtime

- ❖ trading (HTF)



// NYSE trading floor.
// was closed for humans in 2007

- ❖ video / audio processing:



- ❖ orchestration infrastructure:



softrealtime: why JVM (?)



- ❖ Balance (99 % - 'Usual code', 1% - soft realtime)



- ❖ crossing JVM boundaries is expensive.
- ❖ don't do this. [if you can]

Softrealtime & JVM: Programming techniques.


- ❖ Don't guess, know.
- ❖ Measure
 - ❖ profiling [not alw. practical];
 - ❖ sampling (jvm option: -hprof) [jvisualvm, jconsole, flight recorder]
 - ❖ log safepoint / gc pauses
 - ❖ Experiments
 - ❖ benchmarks: <http://openjdk.java.net/projects/code-tools/jmh/>
 - ❖ critical patch: N2N flow.
- ❖ Know details, how code is executed.

Programming techniques: Concurrent models

- ❖ Concurrent flows:

- ❖ minimise context switches:

- ❖ don't switch [number of flows < number of all processor cores].




- ❖ Java: main-thread + service-thread + listener-thread; [user threads: cores-3]

- ❖ pin thread to core [thread affinity]

- ❖ JNI [JNA], exists OSS libraries [OpenHFT]

Programming techniques: data exchange

- ❖ Concurrent flows:
 - ❖ data exchange between threads:
 - ❖ each thread have own ‘local’ memory state.



// state can be unsynchronised
// local access is much faster than volatile
L1 cache: 0.3ns
L2 cache: 1ns
L3 cache: 10ns
RAM: 120ns

Volatile write = RAM write + invalidate cache lines.

- ❖ queue between threads is faster then shared access.
- ❖ locks are evil.

Programming techniques: memory issues

- ❖ GC
 - ❖ young generation (relative short)
 - ❖ full.
 - ❖ [Stop-the-world], pause ~ heap size
- ❖ Contention
 - ❖ flushing cache lines is expensive

Soft-realtime: memory

- ❖ Object: what is java object for the JVM ?

```
class Order
{
    long      id
    int       timestamp;
    String    symbol;
    long      price;
    int       quantity;
    Side      side;
    boolean   limit;
    Party     party;
    Venue     venue;
}
```

1002,t,"APPL",p,q,Bu
y/Sell, Owner,
Venue.

Soft-realtime: memory

- ❖ Object: what is java object for the JVM ?

```
class Order
{
    long      id;
    int       timestamp;
    String    symbol;
    long      price;
    int       quantity;
    Side     side;
    boolean   limit;
    Party    party;
    Venue    venue;
}
```

t, "APPL", p, q, Buy/
Sell, Owner, Venue.

header [64]:
class [32|64]
id [64]:
timestamp 32:
symbol [32-64];
price [64]
quantity[32]
side [32|64]
limit [32]
party [32|64]
venue [32|64]
... padding [0-63]

header, class: 128
data: 64 "APPL"

header, class: 128
.....

header, class: 128
.....

pointer (direct or compressed)

direct = address in memory

Soft-realtime: memory

- ❖ Object: what JMV do ?

```
Order order = new Order()           // alloc 80 bytes from heap  
                                     // order in a local variables (GC root)  
                                     // can force GC if memory reached threshold.
```

```
int t = order.getTimestamp();       // memory read
```

```
Party p = order.getParty();        // memory read
```

```
order.setParty(p1);               // memory write +  
                                     // mark object to be rechecked in next GC cycle
```

Soft-realtime: memory

- ❖ How ‘idiomatic’ Java code will behave ?

```
List<Order> send = new ArrayList<>();  
while(!buffer.isEmpty()) {  
    Order order = buffer.readOrder();  
    int fullPrice = order.price*order.quantity;  
    if (order.getParty().getLimit()+fullPrice > maxLimit) {  
        client.rejected(order,"party-limit")  
    }else{  
        engine.process(order);  
        send.add(order);  
    }  
}
```

Soft-realtime: memory

- ❖ How ‘idiomatic’ Java code will behave ?

```
List<Order> send = new ArrayList<>();
while(!buffer.isEmpty()) {
    Order order = buffer.readOrder();
    int fullPr - order will be created tity;
    if (order.getQuantity() * order.getPrice() > maxLimit) {
        client.rejected(order, "party-limit")
    }else{
        engine.process(order);
        send.add(order);
    }
}
```

- added to long-lived collections

will cause ‘StopTheWorld’

Soft-realtime: memory

- ❖ V1: Die young or live forever.

```
LongSet send = LongSet.create();
while(!buffer.isEmpty()) {
    Order order = buffer.readOrder();
    int fullPr - order will be created tity;
    if (order.getParty() > maxParty) {
        client.rejected(order, "party-limit")
    }else{
        engine.process(order);
        send.add(order.getId());
    }
}
```

- collected in young generation

Soft-realtime: memory

- ❖ V2: Preallocate all.

- order will be created

```
LongSet send = LongSet.create();
Order order = new Order();
while(!buffer.isEmpty()) {
    buffer.fillOrder(order);
    int fullPrice = order.price*order.quantity;
    if (order.getParty().getLimit()+fullPrice > maxLimit) {
        client.rejected(order,"party-limit")
    }else{
        engine.process(order);
        send.add(order.getId());
    }
}
```

- one ‘static’ object per thread
 - no allocations.

Soft-realtime: memory

- ❖ V3: Access to serialised data instead object

```
class OrderAccess
{
    static void getId(Buffer b,int offset)
    { return b.getLong(offset+0) }

    static void putId(Buffer b, int offset, long id)
    { b.putLong(offset+0,id) }

    static int getTimestamp(Buffer b, int offset)
    { return b.getInt(offset+8); }

    .....

}
```

Soft-realtime: memory

- ❖ V3: Access to serialised data instead object

```
LongSet send = LongSet.create();
for(int offset=0; !buffer.atEnd(offset);
                offset+=OrderAccess.SIZE) {
    long price = OrderAccess.getPrice(buffer,offset);
    int quantity = OrderAccess.getQuantity(buffer,offset);
    int fullPrice = price * quantity;
    long partyLimit = OrderAccess.getPartyLimit(buffer,offset);
    if (partyLimit > maxLimit) {
        client.rejected(buffer,offset,"party-limit")
    }else{
        engine.process(buffer,offset);
        send.add(order.getId());
    }
}
```

- no temporary objects
- no allocations
- slightly unusual API

Soft-realtime: memory

- ❖ V4: Offheap memory access via unsafe

sun.misc.Unsafe

- ❖ internal low-level sun API
- ❖ implemented by JIT intrinsic
- ❖ JVM crash on error
- ❖ widely used but not recommended
- ❖ will be deprecated in Java-9
 - ❖ (removed in Java 10)
 - ❖ (we will discuss alternatives)

```
class Unsafe
{
    byte getByte(long addr)
    void setByte(long addr, byte value)
    ...
    byte getByte(Object value, long offs)
    ...
    // same for primitive types.
    ...
    // access to Object internals.
}
```

Soft-realtime: memory

- ❖ V4: Offheap memory access via unsafe

```
class OrderAccess
{
    static void getId(long address)
    {   return unsafe.getLong(address) }

    static void putId(long address, long id)
    {   unsafe.putLong(offset+0,id) }

    static int getTimestamp(Buffer b, int offset)
    {   return b.getInt(offset+8); }

    .....

}
```

Soft-realtime: memory

❖ V4: Offheap memory access via unsafe

```
LongSet send = LongSet.create();
for(long address=buffer.begin; address!=buffer.end();
    address=advance(address,OrderAccess.SIZE) ) {
    long price = OrderAccess.getPrice(address);
    int quantity = OrderAccess.getQuantity(address);
    int fullPrice = price * quantity;
    long partyLimit = OrderAccess.getPartyLimit(address)
    if (partyLimit > maxLimit) {
        client.rejected(address,"party-limit")
    }else{
        engine.process(address);
        send.add(order.getId());
    }
}
// better wrap by more hight-level interfaces (buffer, etc)
```

- no allocations
- any memory
- subtle bugs are possible.

Soft-realtime: memory

- ❖ V3, V4 Problem: Verbose API
 - bytecode transformation from more convenient API
 - Metaprogramming within hight-level JVM languages.
<http://scala-miniboxing.org/ildl/>
<https://github.com/densh/scala-offheap>
 - Integration with unmanaged languages
 - JVM with time-limited GC pauses. [Azul Systems]

Memory access API in next JVM versions

- ❖ VarHandlers (like methodHandlers)
 - ❖ inject custom logic for variable access
 - ❖ JEP 193 <http://openjdk.java.net/jeps/193>
- ❖ Value objects.
 - ❖ can be located on stack
 - ❖ JEP 169. <http://openjdk.java.net/jeps/169>
- ❖ Own intrinsics [project Panama]
 - ❖ insert assembler instructions into JIT code.
 - ❖ <http://openjdk.java.net/projects/panama/>
 - ❖ JEP 191. <http://openjdk.java.net/jeps/191>
- ❖ ObjectLayout [Arrays 2.0]
 - ❖ better contention
 - ❖ <http://objectlayout.github.io/ObjectLayout/>

Low latency Java: access to native code

- ❖ JNA
 - ❖ (idea — API to build native stack call and get result)
 - ❖ pro: not require packaging native libraries.
 - ❖ cons: SLOWWWW (but one-time: ok)
- ❖ JNI
 - ❖ (idea — bridge between JNI/native code)
 - ❖ pro: faster than JNA, exists critical jni
 - ❖ cons: need packaging native library
- ❖ Shares Memory
 - ❖ (idea — use same binary address from native & jvm code)
 - ❖ pro: fast
 - ❖ cons: unsafe, unusual API

Low latency Java: JNI

- ❖ JNI
 - ❖ (idea — bridge between JNI/native code)
 - ❖ Java Object => Handle for native (maintain table)

A Sample Native Call

```
save 64      # register window push
add g7,&jnienv_offset,o0 # JNIEnv* in o0
set 0,o1      # assume null arg
beq i0,skip   # handle of null is null
stw i0,[sp+72] # handleize 'this'
add sp,72,o1   # ptr to 'this' passed in o1
skip:
std d0,[sp+64] # float args passed in int regs
ldd [sp+64],o2 # double now in o2/o3
setlo #ret_pc,10
sethi #ret_pc,10 # set 32-bit PC before allow GC
stw 10,[g7+&pc] # Rely on Sparc TSO here (IA64!)
stw sp,[g7+&sp] # Enable GC from here on
call foo      # The Native Call
ret_pc: add g7+4sp,10 # CAS does not take an offset
cas 0,sp,[10]  # if SP still there, store 0 to disable gc
hne gc_in_progress
set 0,i0      # assume null result
beq o0,is_null
ldw [o0],i0    # de-handleize return value
is_null:
return
restore      # pop sparc window
```

Handles,
GC-unsafe natives

Arg-shuffle
JIT vs native convention

Stack crawl
GC lock

The Call

GC unlock

De-Handleize



// part of Cliff Click presentation
(advise for VM writers)

Low latency Java: JNI

- ❖ JNI

- ❖ (idea — bridge between Java and native)
- ❖ Java Object => Handle

```
class JNIWrapper
{
    void native eventLoop();
    void processCallback(Packet p);
}
```

- ❖ Callbacks are extremely slow

```
C:
extern Java_nativeEventLoop(JNIEnv jenv, jobject job) {
    while(!endOfBuffer(buff)) {
        packet p = getData(buff);
        if (checked(p)) {
            processInJavaCallback(p); - slow
        }
    }
}
- p will be marshalled / unmarshalled each time
}
```

Low latency Java: JNI

// better use return instead callback

```
extern Java_nativeEvent(JNIEnv jenv, jobject job) {  
    static int state = START;  
    switch(state) {  
        case START:  
            state = LOOP;  
            while(! endOfBuffer()) {  
        case LOOP:  
            Packet p = readData();  
            if (checked(p)) {  
                return (long)p;  
            }  
        }  
    }  
    return -1;  
}
```

```
class JNIWrapper  
{  
    void native event();  
    void process(long packetAddr);  
}  
  
...  
  
for(;;) {  
    p=event();  
    if (p== -1L) break;  
    process(p);  
}
```

Low latency Java: JNI

- ❖ Critical JNI
 - ❖ not in official documentation
 - ❖ substituted instead JNI after JIT warm-up.
 - ❖ only primitive types; static calls.
 - ❖ must not be long-running. (GC is disabled)


```
extern long Java_nativeEvent(JNIEnv jenv, jobject job)
->
    extern long JavaCritical_nativeEvent()
```

// want add timeout to satisfy 'last' condition

Low latency java: shared memory interface

- ❖ idea:
 - ❖ spawn process or native thread.
 - ❖ mmap file with given name to memory on ‘native side’.
 - ❖ mmap file with given name to memory on ‘java’ side.
 - ❖ have shared memory persist to the same file in java & native code.

Low latency Java/Native code data exchange



Low latency Java/fw

- ❖ Low-latency on JVM is
 - ❖ possible
 - ❖ sometimes non-trivial
 - ❖ special techniques exists, in future will be better.



Thanks for attention.

❖ Questions ?

// ruslan@shevchenko.kiev.ua

// [@rssh1](https://github.com/rssh)

// <https://github.com/rssh>