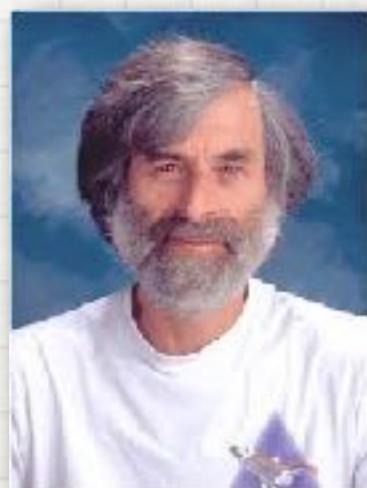


PAXOS (AND A LITTLE ABOUT OTHER CONSENSUSES)

//RUSLAN SHEVCHENKO



Leslie Lamport
The Part-Time Parliament



PAXOS PAPERS

1. Leslie Lamport. The Part-Time Parliament.

ACM Transactions on Computer Systems. 1990 — rejected
resubmitted in 1998 — published

<https://www.microsoft.com/en-us/research/publication/part-time-parliament/>

2. Leslie Lamport. Paxos Made Simple

ACM SIGACT News (Distributed Computing Column) 32, 4
(Whole Number 121, December 2001) 51-58.

<https://lamport.azurewebsites.net/pubs/paxos-simple.pdf>

3. Tushar Deepak Chandra, Robert Griesemer, Joshua Redstone

Paxos Made Live - an Engineering Perspective.

ACM, Symposium on Principles of Distributed Computing. 2007

<https://ai.google/research/pubs/pub33002>

NON-PAXOS PAPERS

// will talk a little, details can be a theme for another PWL.

1. Diego Ongaro and John Ousterhout

In Search of an Understandable Consensus Algorithm

2014 USENIX Annual Technical Conference.

<https://www.usenix.org/system/files/conference/atc14/atc14-paper-ongaro.pdf>

<https://raft.github.io/>

2. Miguel Castro and Barbara Liskov

Practical Byzantine Fault Tolerance

Third Symposium on Operating Systems Design and Implementation,

New Orleans, USA, February 1999

<http://pmg.csail.mit.edu/papers/osdi99.pdf>

<https://github.com/luckydonald/pbft>

3. Leslie Lamport.

Byzantizing Paxos by Refinement.

Distributed Computing: 25th International Symposium: DISC 2011, David Peleg, editor. Springer-Verlag (2011) 211-224.

<http://lamport.azurewebsites.net/pubs/web-byzpaxos.pdf>

DISTRIBUTED SYSTEMS: FUNDAMENTALS

- N nodes.



- Can send messages via network

- $\text{Send}(m, i, j)$

We don't know: is it delivered

NETWORK

ASYNCHRONOUS

- $\text{Receive: } j \Rightarrow (m, i)$

know, when node or channel is down

ASYNCHRONOUS
WITH FAILURE DETECTOR

ping/pong timeout

universal time.

PARTIALLY
SYNCHRONOUS

i

j

$\text{send(ping, } i, j)$

$\text{receive(ping, } i, j)$

$\text{send(pong, } j, i)$



SYNCHRONOUS

$\text{receive(pong, } j, i)$

$\exists \delta t : t_i(\text{send(ping, } i, j)) - t_i(\text{receive(pong, } i)) < \delta_t$

DISTRIBUTED SYSTEM: FUNDAMENTALS

CONSENSUS PROBLEM



*“We have an agreement in principle.
The question is, do we all have the same principles?”*

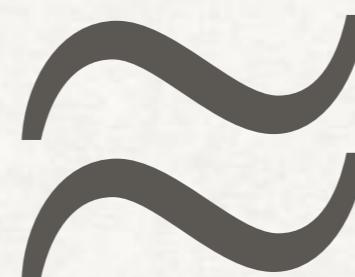
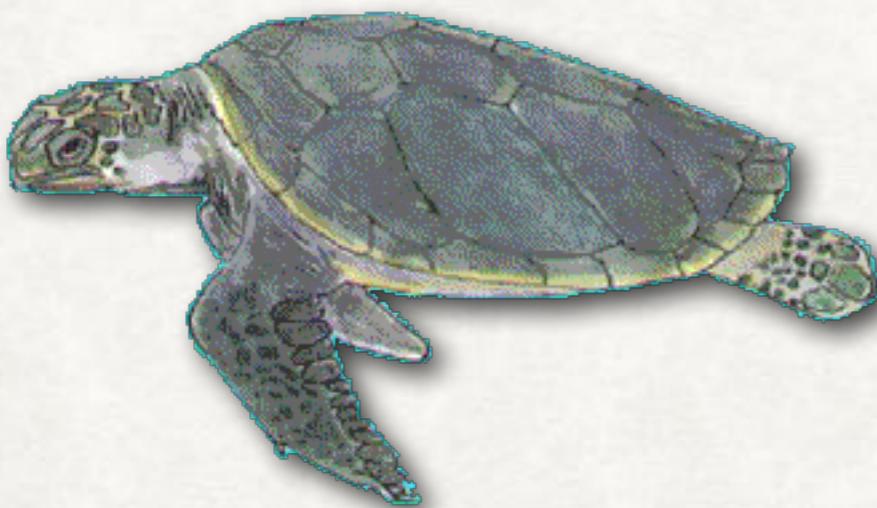
p_i - program on node i.

$\exists c : \forall i, j \in N : result(p_i) = c$

DISTRIBUTED SYSTEMS: FUNDAMENTALS

ASYNCHRONOUS CONSENSUS IS IMPOSSIBLE

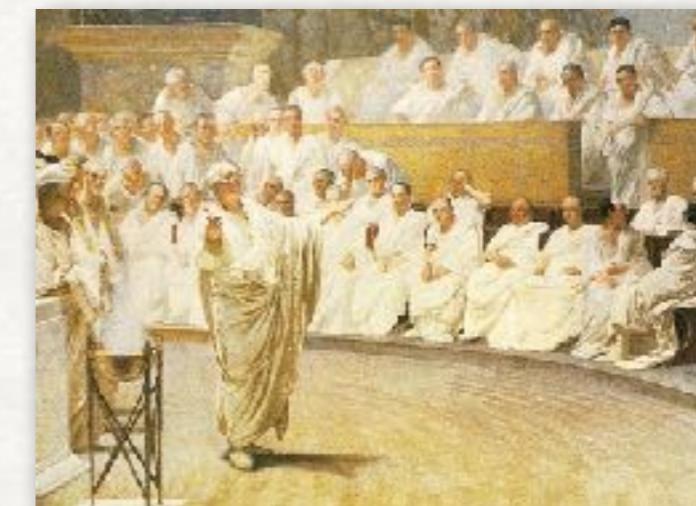
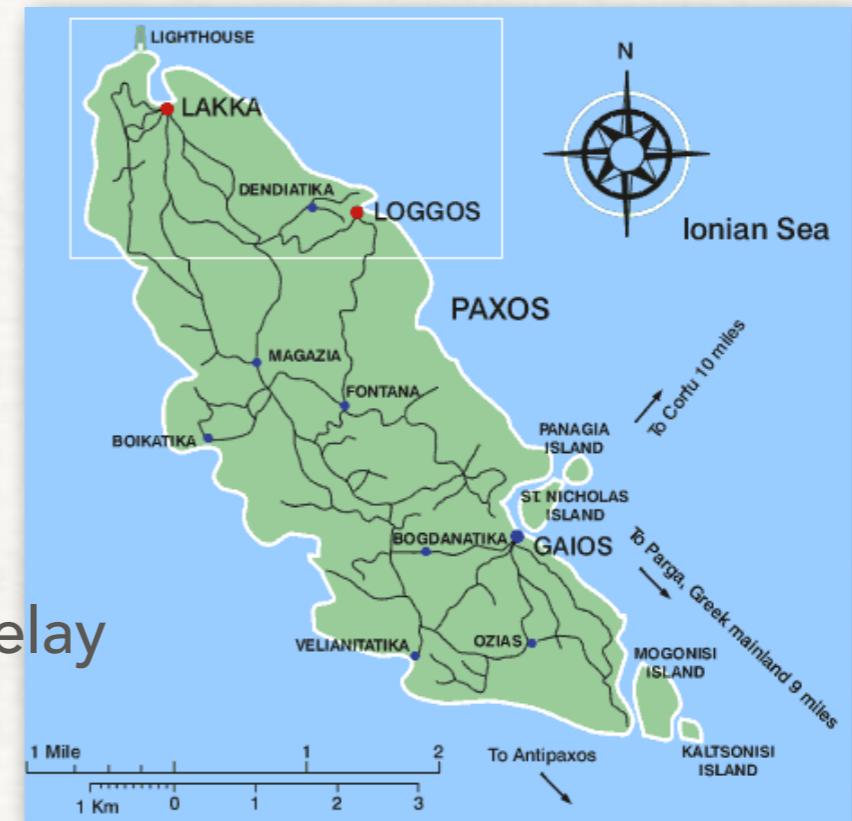
- Fischer, Lynch and Patterson. 'Impossibility of Distributed Consensus with One Faulty Process', J. ACM 32, 2 (April 1985), 374-382. <http://cs-www.cs.yale.edu/homes/arvind/cs425/doc/fischer.pdf>
- Proof Idea:



PAXOS

PART-TIME PARLIAMENT

- Archeologist, describing voting system.
- Parliament 'work from home/travel'
- Messages are sent by part-time courier
- Any member at any time can fire issue
- When the majority of votes is collected,
the decision had to be done
- members can be non-reliable or reply with long delay



PACTS

PART-TIME PARLIAMENT

Roles:

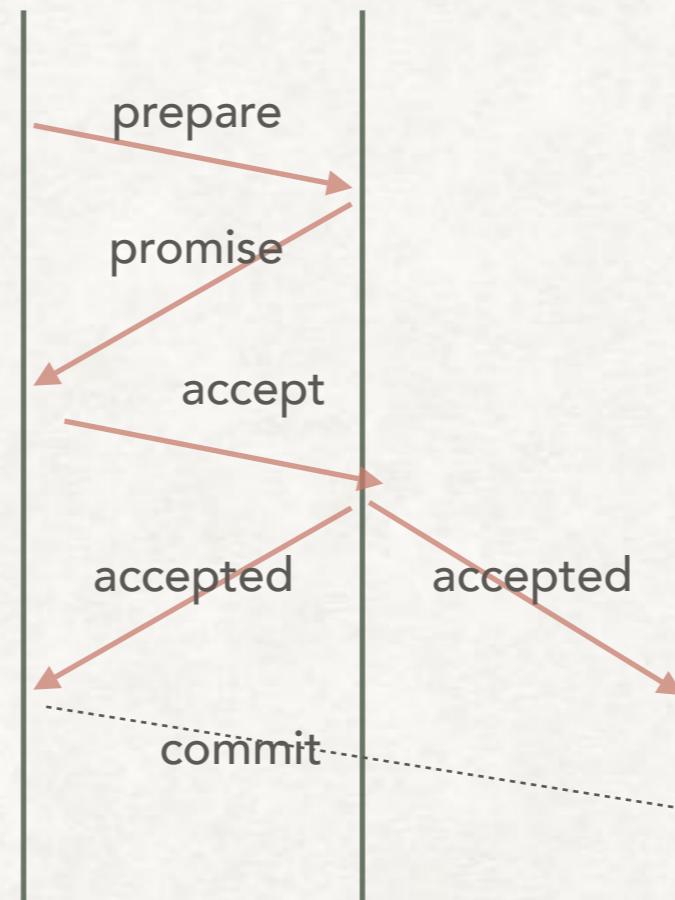
- Proposer,
- Acceptor
- Learner

On practice, merged in one.

Messages:

- prepare
- promise
- accept
- accepted
- optional, implicit

Proposer Acceptor Learner



PACTS

PART-TIME PARLIAMENT

Roles:

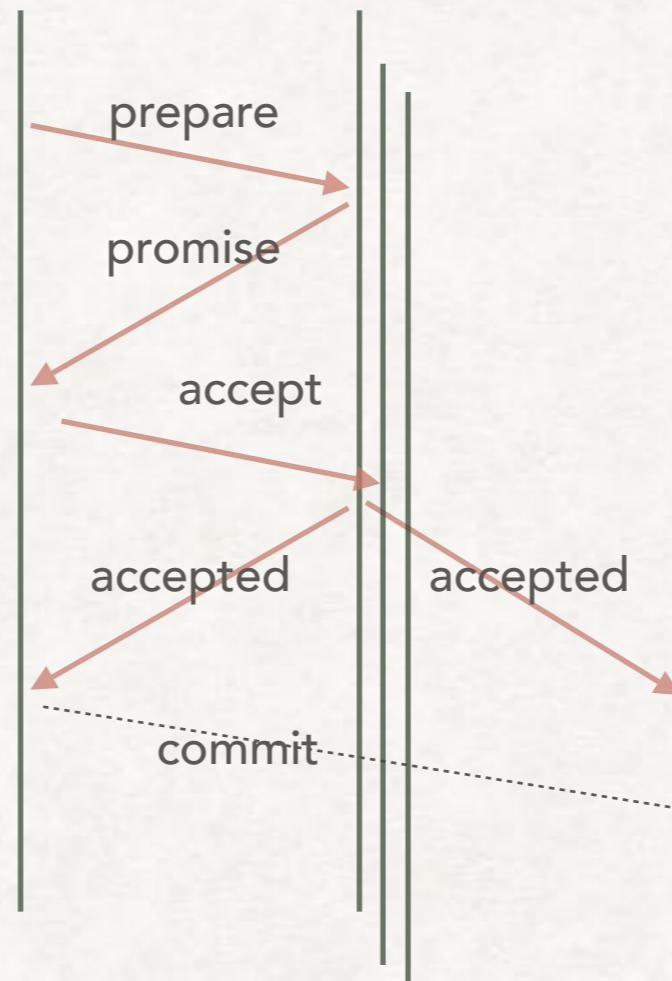
- Proposer,
- Acceptor
- Learner

On practice, merged in one.

Messages:

- prepare
- promise
- accept
- accepted
- optional, implicit

Proposer Acceptors Learner



Quorum. $|i \in Q| > N/2$

PACTOS FORMALIZATION

- **Nodes:** $n_i \in N$, each node have number
 $issue \in I$
 $value \in V$

$prepare(n_{proposer} : N, issue : I, round : N)$

$promise(n_{acceptor} : N, issue : I, round : N, value : Option[V])$

$accept(n_{proposer} : N, issue : I, round : N, value : V)$

Optional

$accepted(n_{acceptor} : N, issue : I, round : N, value : V)$

$acceptNack(n_{acceptor} : N, issue : I, round : N)$

$ballot(message) = (message.value, message.round)$

PAXOS

PROPOSER

```
State: myProposal: Option[Value], round:Nat,    prepare( $n_{proposer} : N, issue : I, round : N$ )
round <- myN
broadcast(Prepare(me, issue, round))

receive{
  case p:Promise if checkRound(p) =>
    myProposal=best(myProposal, ballot(p))
    if (prepare-quorum-ready) {
      value = myProposal.getOrElse(random)
      broadcast Accept(me, issue, round, value)
    }
    if (quorum-failed) {
      fail-round
    }
  case m: Accepted if (checkRound(m)) =>
    if (accepted-quorum-ready(m)){
      //local Commit(myProposal, round)
      finish(myProposal)
    } else if (quorum-failed) {
      failRound
    }
  case m:Commit => finish(m.value)
}

def fail-round() = {
  round = ((round/N)+1)*N + myN
  myProposal = None
  wait-some-time
  restart()
}
```

PAXOS

ACCEPTOR

```
State: myValue: Option[Value], round:N,  
  
receive{  
  case p:Prepare =>  
    r = if (myValue.isEmpty) p.round  
         else this.round  
    reply Promise(me,p.issue,r,myValue)  
  
  case m: Accept if (p.round > round) =>  
    myValue match {  
      case None =>  
        myValue = p.value  
        round = p.round  
        reply Accepted(me,m.issue,round,m.value)  
      case Some(v) =>  
        if (m.value == v) {  
          reply Accepted(me,m.issue,round,m.value)  
        } else {  
          reply AcceptedNack  
        }  
    }  
  case m:Commit => finish(m.value)  
}
```

$$\text{Accepted}(v, r) \Rightarrow \forall k > r : \text{Promise}(v, k)$$

PAXOS

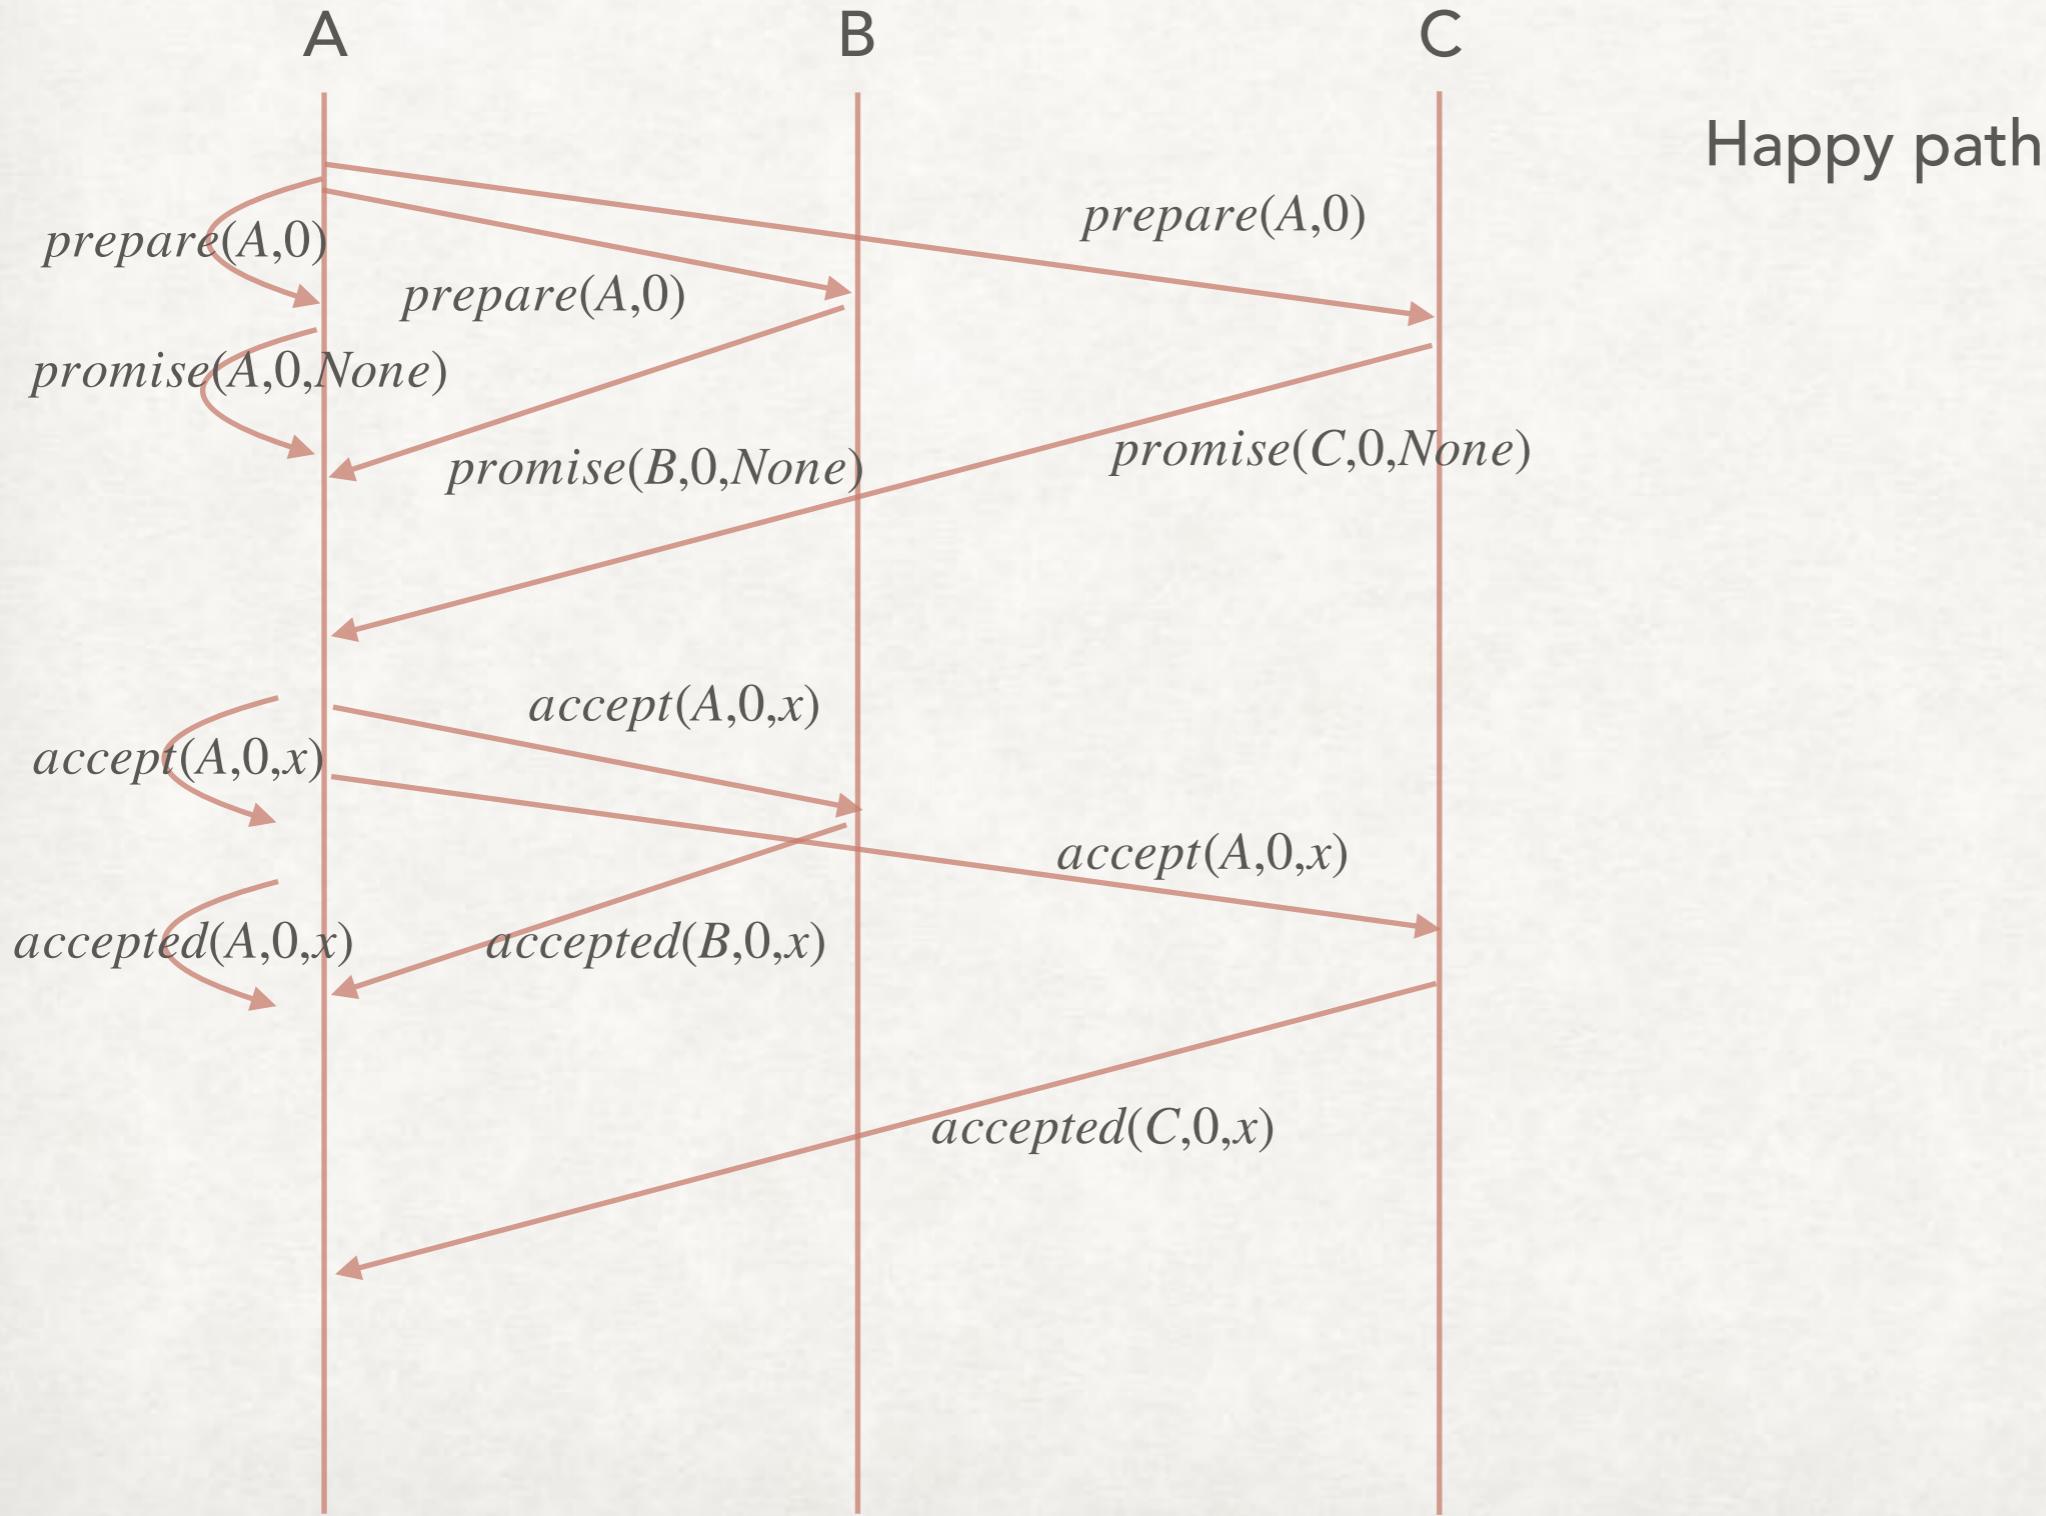
LEARNER

```
State: myValue: Option[Value], round:N,  
  
receive{  
  case m:Accepted =>  
    if (checkQuorum(m)) {  
      learn(m.issue,m.value)  
        // generate local commit(m)  
    }  
}  
}
```

// On practice, Proposal, Acceptor, Learner — the same.

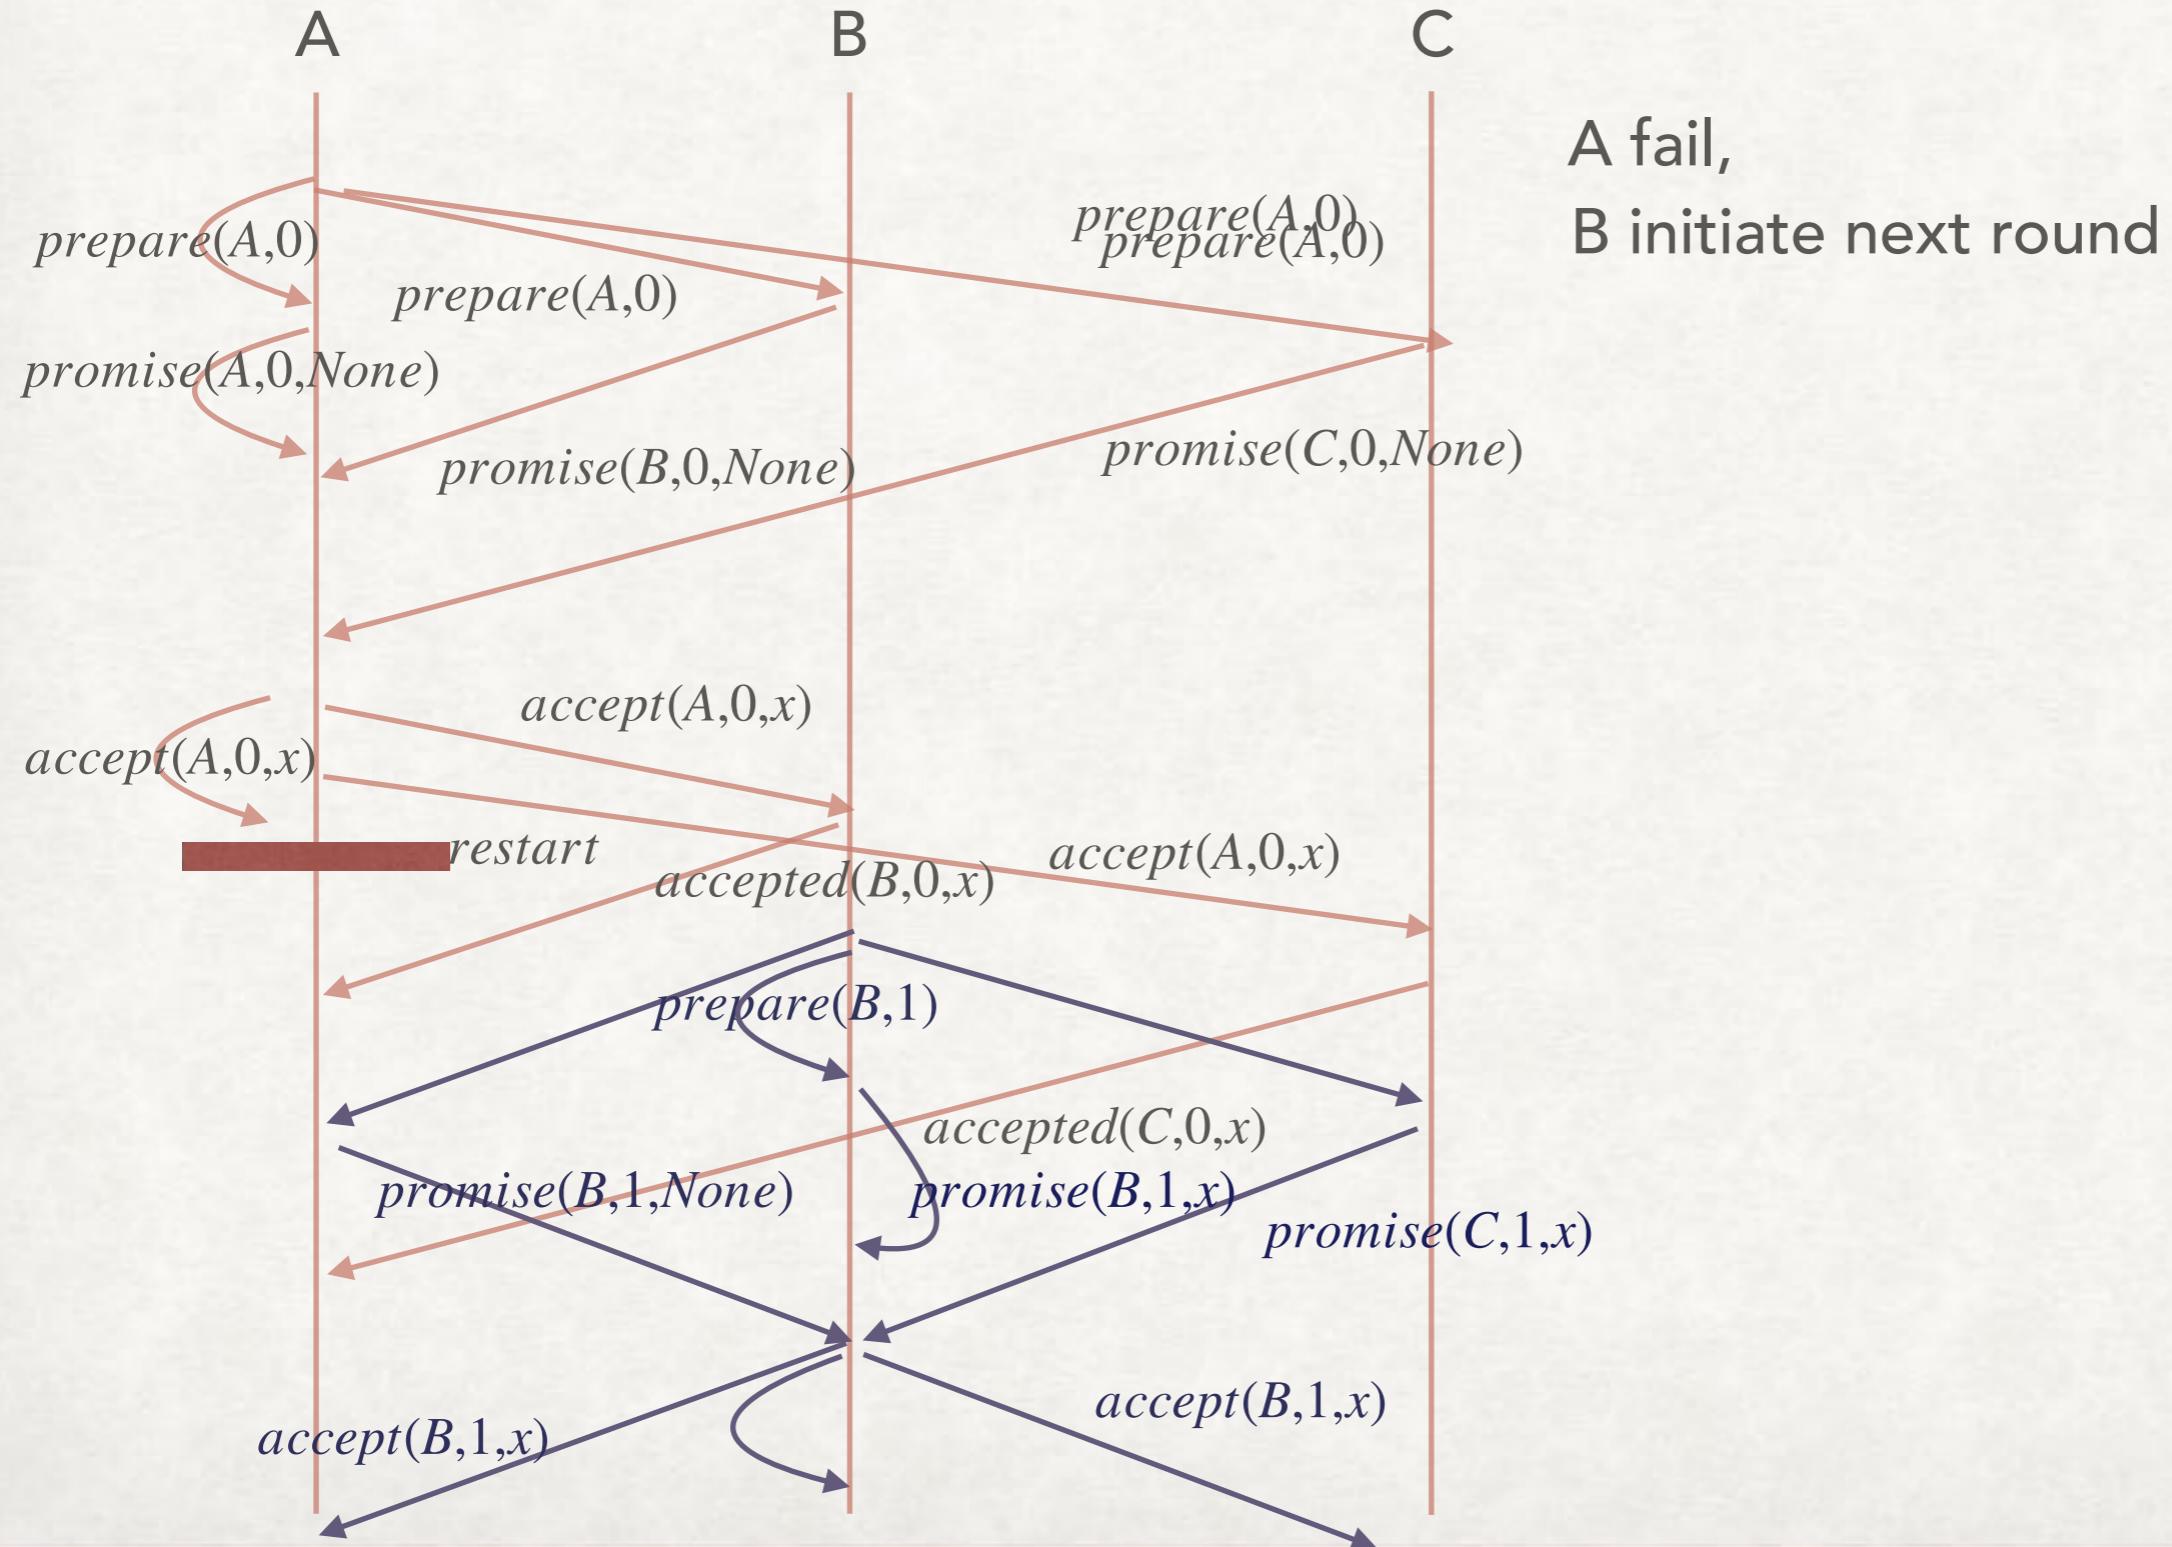
PAXOS - RUN

LET WE HAVE 3 NODES, ALL PLAY BOTH ROLES



PAXOS - RUN

LET WE HAVE 3 NODES, ALL PLAY BOTH ROLES



PROPERTIES

PROOFS

- Safety:
 - Only a value that has been proposed may be chosen.
 - Only a single value is chosen. (non-trivial)
 - Only chosen values are learned.
- Liveness impossible (due to FLP), eventual progress instead:
 - Proposal will learn previous highest number

PAXOS SAFETY

ONLY A SINGLE VALUE HAS CHOSEN

$\forall n \in N : Accepted(n, round, v) \rightarrow \forall k > round : Accepted(n, k, v)$

$\exists Q \in Quorum(N) : \forall n \in Q : Accepted(n, r, v) \rightarrow Chosen(v, r)$

$Chosen(v, r) \rightarrow \forall k > r : Chosen(v, k)$

Let we have 2 values chosen:

$Chosen(v1, r) \& Chosen(v2, k) \rightarrow Chosen(v1, max(r, k)) \wedge Chosen(v2, max(r, k))$

$\exists Q_1 \in Quorum(N), Q_2 \in Quorum(N) : \forall n_1 \in Q_1 \forall n_2 \in Q_2 :$
 $Accepted(n_1, v_1, r) \wedge Accepted(n_2, v_2, r)$

$v_1 \neq v_2 \rightarrow Q_1 \cap Q_2 = \emptyset$

But any two Quorums should be intersect

$Q_1 \in Quorum(N), Q_2 \in Quorum(N) \rightarrow (Q_1 \wedge Q_2 \neq \emptyset) \quad !\text{Contradiction!}$

PAXOS SAFETY

ONLY A SINGLE VALUE HAS CHOSEN

$\forall n \in N : Accepted(n, round, v) \rightarrow \forall k > round : Accepted(n, k, v)$

$\exists Q \in Quorum(N) : \forall n \in Q : Accepted(n, r, v) \rightarrow Chosen(v, r)$

$Chosen(v, r) \rightarrow \forall k > r : Chosen(v, k)$

Let we have 2 values chosen:

$Chosen(v1, r) \& Chosen(v2, k) \rightarrow Chosen(v1, max(r, k)) \wedge Chosen(v2, max(r, k))$

$\exists Q_1 \in Quorum(N), Q_2 \in Quorum(N) : \forall n_1 \in Q_1 \forall n_2 \in Q_2 :$
 $Accepted(n_1, v_1, r) \wedge Accepted(n_2, v_2, r)$

$v_1 \neq v_2 \rightarrow Q_1 \cap Q_2 = \emptyset$

But any two Quorums should be intersect

$Q_1 \in Quorum(N), Q_2 \in Quorum(N) \rightarrow (Q_1 \wedge Q_2 \neq \emptyset) \quad !\text{Contradiction!}$

PACTS PROGRESS

$fail(n, r) \rightarrow \exists n \in N, r_1 > r : Prepare(n, r_1)$

// scheduling run on randomized timeout

Modifications:

Multi-Paxos: same leader process events until failure

not need in Prepare, until we have no failures.

PACT-OS IMPLEMENTATION

CHUBBY – GOOGLE LOCK SERVICE

3. Tushar Deepak Chandra, Robert Griesemer, Joshua Redstone
Paxos Made Live - an Engineering Perspective.

ACM, Symposium on Principles of Distributed Computing. 2007

<https://ai.google/research/pubs/pub33002>

2 pages pseudocode => $k \cdot 10^2$ LOC in C++.

Code generated from own specification language

+ Deterministics variant for testing.

Performance superior to 3DB

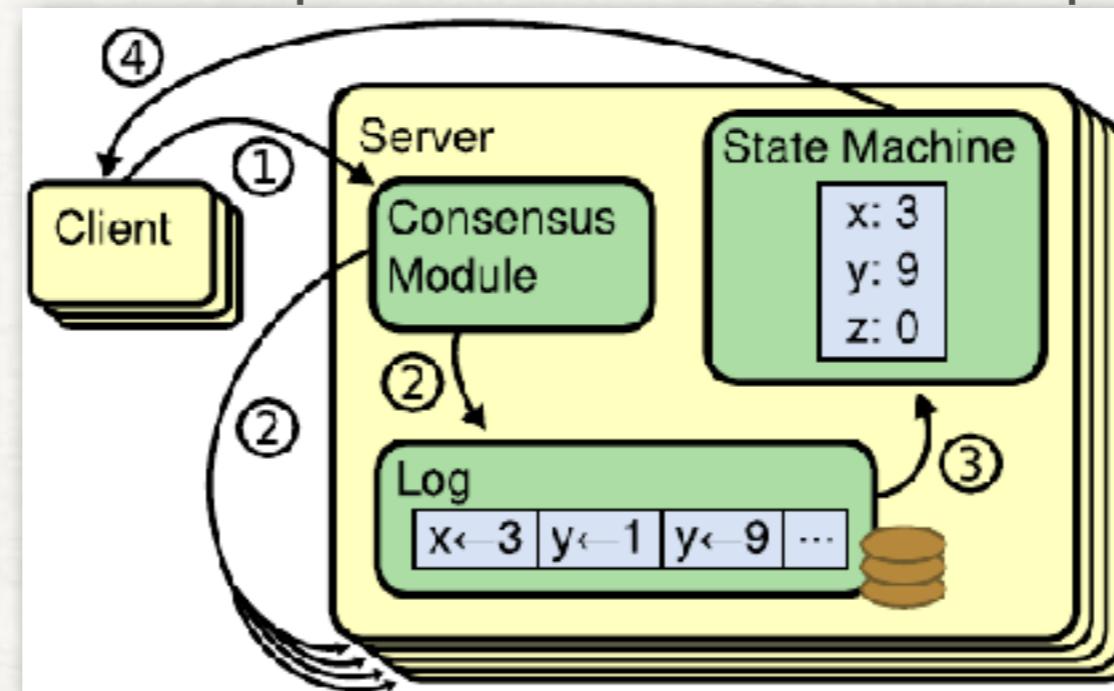
MultiPaxos + optimizations
- lagging.

Testing fail-over solutions is hard.

RAFT

IN SEARCH FOR UNDESTANDABLE CONSENSUS ALGORITHM (2014)

- MultiPaxos usually used as part of state-machine replication.

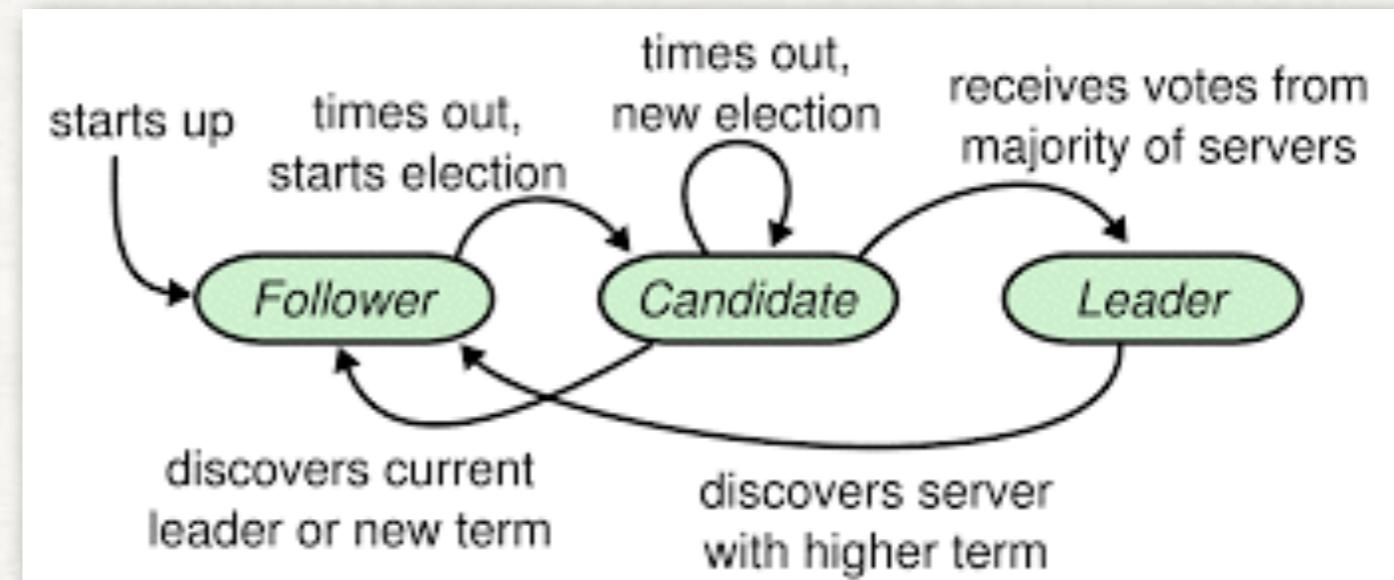


- One writer (leader), do something and send to followers. When leader failed, new leader is elected.
- RAFT - state machine replication, simpler (in theory) than paxos with similar performance characteristics.

RAFT

STATE MACHINE REPPLICATION

- <https://raft.github.io/>
- Instead

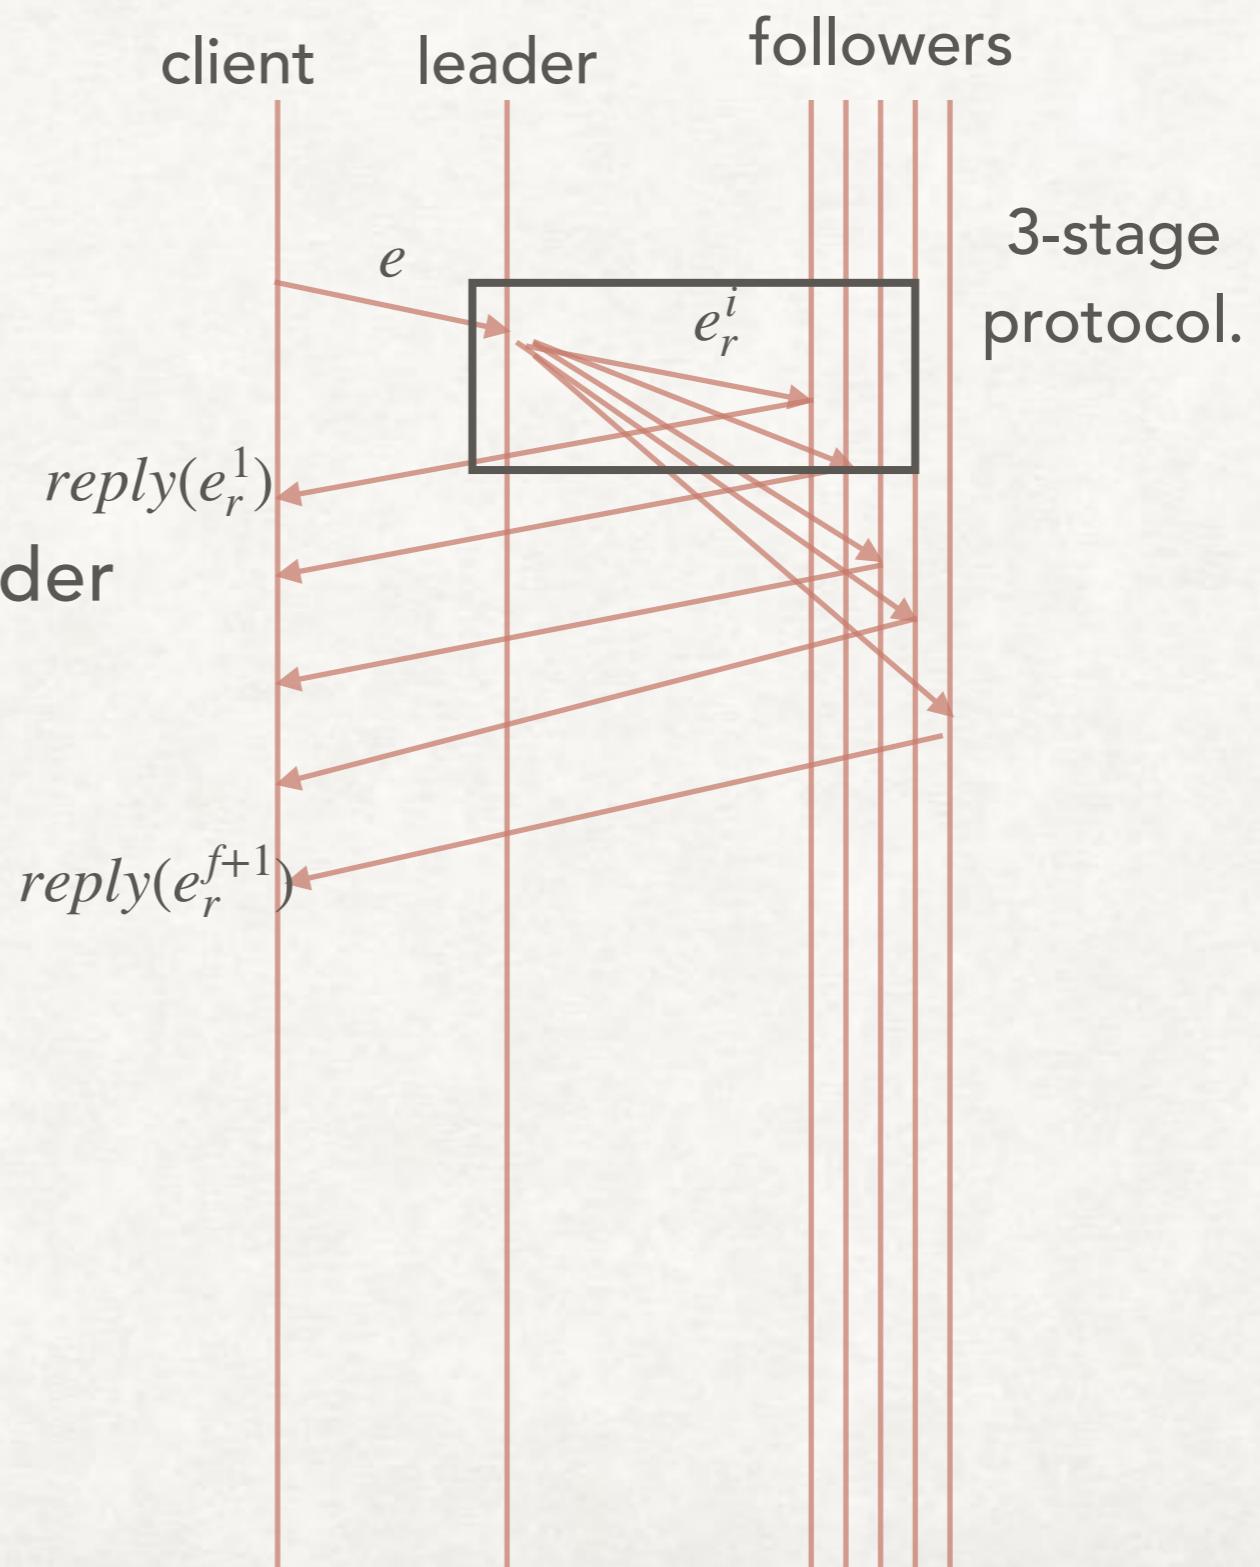


- Can be viewed as 'refactoring' of PaxOS: merging roles into one,
- use heartbeat for failure detection
- on leader failure, node become a candidate and request other node to vote (choose one) as leader.

BPFT

BYZANTINE FAULT TOLERANCE

- Replicated state machine
- Some nodes can lie (f):
 - Client submit item to leader
 - Gather $f+1$ replies.
 - If all replies
 - are the same: all ok
 - otherwise - resend.



BPFT

BYZANTINE FAULT TOLERANCE

- Replicated state machine

- Some nodes can lie (f):

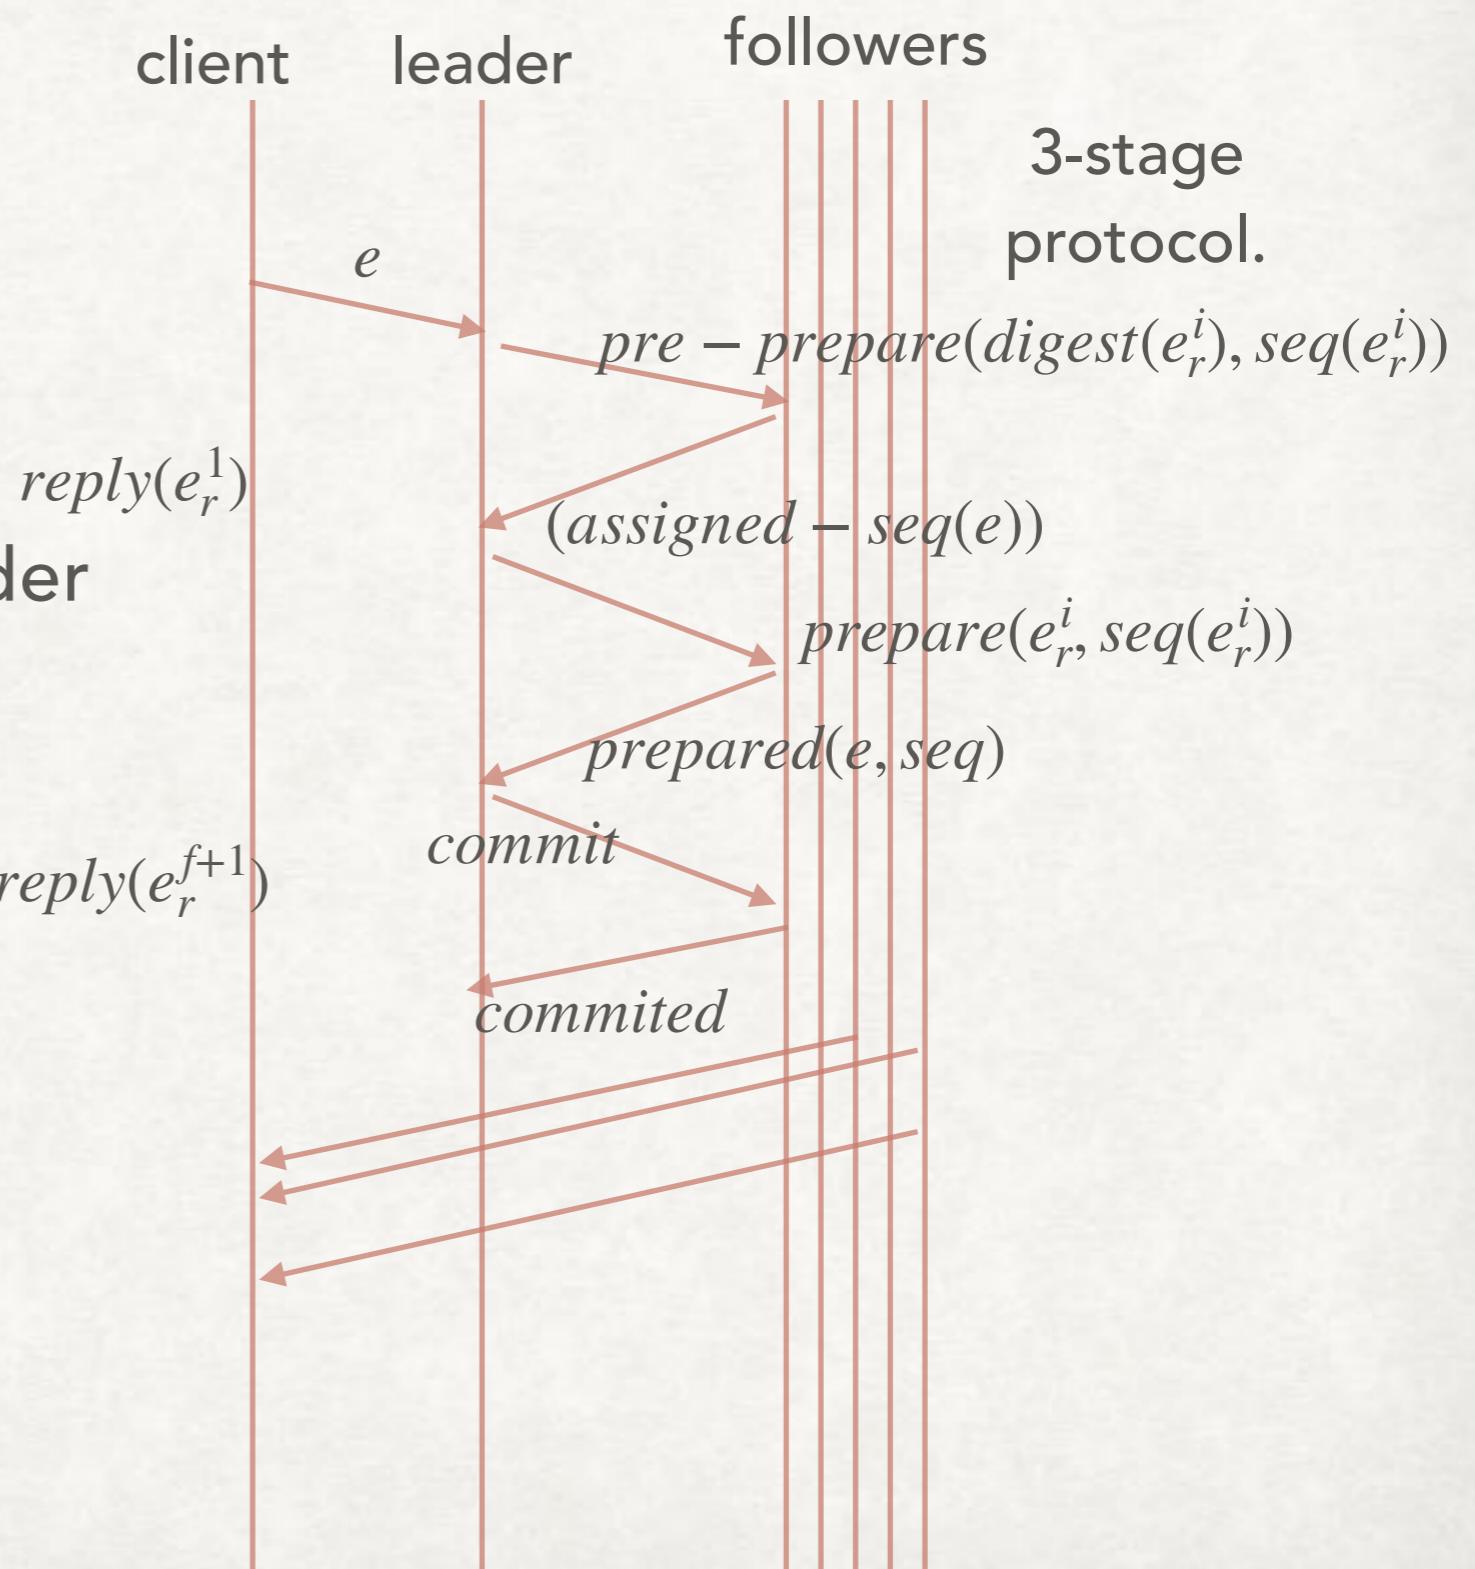
- Client submit item to leader

- Gather $f+1$ replies.

- If all replies

- are the same: all ok

- otherwise - resend.



BYZANTIZING PAXOS BY REFINEMENT

LESLI LAMPORT, 2011

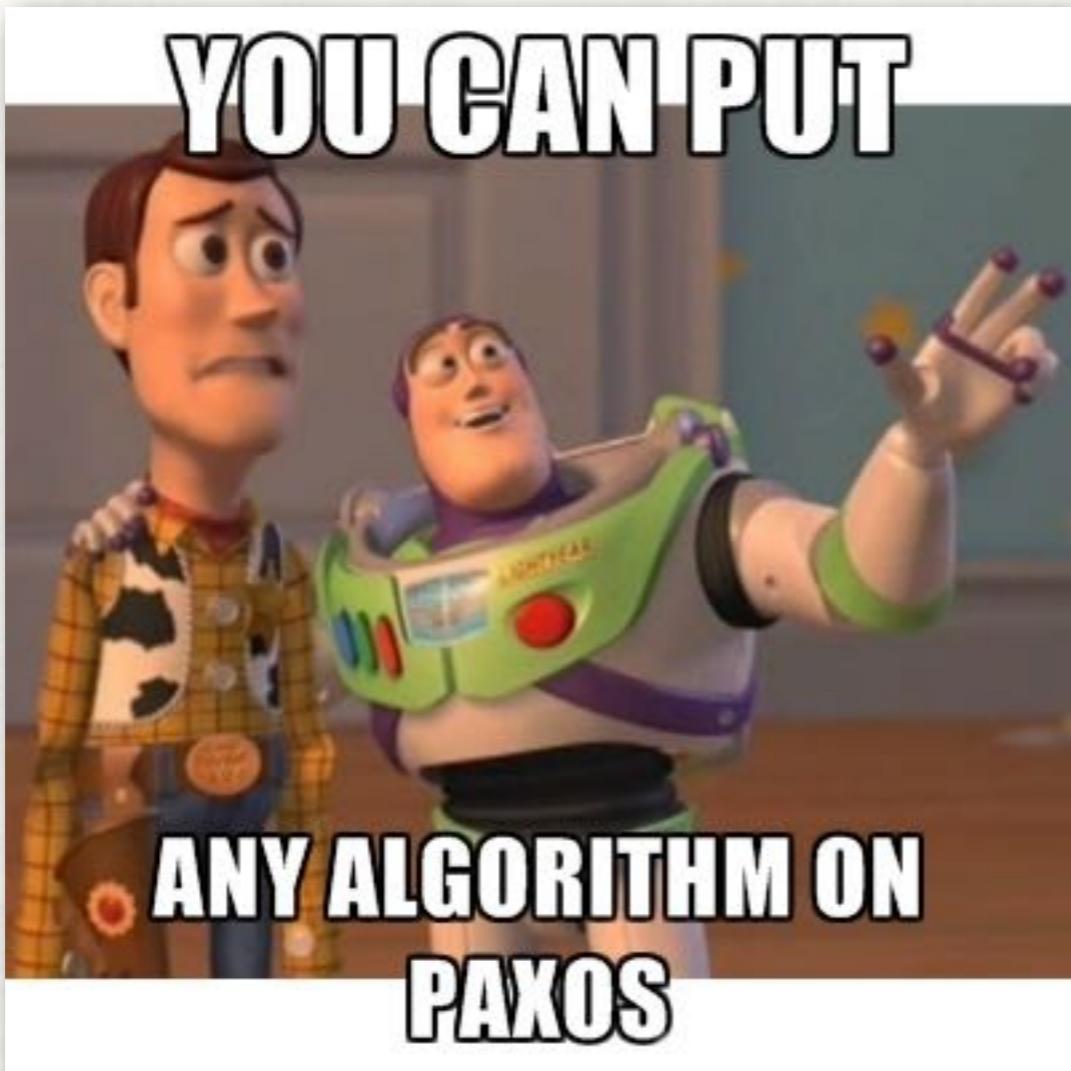
- BPCon — Byzantine fault tolerance (prev. slides)
- PCon — Paxos
- Byzantine: $\text{Algorithm}(N) \Rightarrow \text{Algorithm}(N+f)$
- $\text{Byzantine}(PCon) = BPCon \quad N > 3f$
 $f+1$ round (in worst case)

Idea: $S \rightarrow M \Rightarrow BzQourum(S) \rightarrow M$

$S \rightarrow M \Rightarrow BzQourum(S) \rightarrow \text{Proof}(BzQourum(S), M)$

THANKS FOR LISTENING

PAPERS WE LOVE, KYIV, 27.08.2018



//The Story, told by Ruslan Shevchenko.

ruslan@shevchenko.kiev.ua
@rssh1